COMP3711: Design and Analysis of Algorithms

WU Jinjun

October 12, 2026

Contents

1 Prerequisite: Asymptotic Notations **3**

2 Divide and Conquer **4**
2.1 Binary Search . . . 4
2.1.1 Variant: Find the Boundary (lower_bound) . . . 5
2.1.2 Exercise 1: Rotated Sorted Array . . . 6
2.1.3 Exercise 2: Finding the Last 0 . . . 7
2.2 Mergesort & Inversion Count . . . 7
2.2.1 Warm-up: Tower of Hanoi . . . 7
2.2.2 Merge: Combining Two Sorted Halves . . . 8
2.2.3 Mergesort Algorithm . . . 9
2.2.4 Counting Inversions . . . 10
2.3 Maximum Subarray . . . 11
2.3.1 Application: Maximizing Stock Profit . . . 13
2.4 Multiplication and Master Theorem . . . 13
2.4.1 Integer Multiplication (Karatsuba) . . . 13
2.4.2 First D&C Attempt: Four Half-Size Products . . . 13
2.4.3 Master Theorem for Recurrences . . . 15
2.5 Matrix Multiplication and Strassen's Algorithm . . . 16

3 Randomized Algorithms **17**
3.1 Why Randomization? . . . 17
3.2 Probability Review . . . 17
3.3 Linearity of Expectation . . . 17
3.4 Example: Hiring Problem . . . 18
3.5 More Indicator-Variable Examples . . . 18
3.5.1 Guessing Shuffled Cards . . . 18
3.5.2 Hash Table with Chaining . . . 18
3.5.3 Birthday Paradox . . . 19
3.6 Generating a Uniform Random Permutation . . . 19
3.7 Geometric Distribution and Coupon Collector . . . 19
3.8 Case Study: The Majority Element . . . 19

4 Quicksort **21**
4.1 Deterministic Quicksort (Lomuto Partition) . . . 21
4.2 Quicksort Analysis . . . 22
4.3 Randomized Selection (Quickselect) . . . 23
4.4 Hoare's Partition (Alternative) . . . 24

5 Heapsort **25**
5.1 Priority Queue Motivation . . . 25
5.2 Heap Data Structure . . . 25
5.3 Max-Heapify (Bubble-Down) . . . 25
5.4 Build Max-Heap . . . 26
5.5 Heapsort Algorithm . . . 27
5.6 Priority Queue Operations (Applications of Heaps) . . . 27
5.7 Applications . . . 28
5.7.1 Merging k Sorted Arrays . . . 28
5.7.2 Delete and Update-Priority . . . 28
5.7.3 The Skyline Problem . . . 28

6 Linear-Time Sorting **29**
6.1 The Decision-Tree Model and the Sorting Lower Bound . . . 29
6.2 Counting Sort . . . 30
6.3 Radix Sort . . . 30
6.4 Bucket Sort . . . 31

7 Greedy Algorithms **33**
7.1 Warm-up: Change-Making . . . 33
7.2 Warm-up: Max Sum of Non-Adjacent Elements . . . 33
7.3 Interval Scheduling . . . 34
7.3.1 Three Tempting Greedy Rules (All Fail) . . . 34
7.3.2 The Correct Rule: Earliest Finish Time . . . 34
7.3.3 Correctness Proof (Exchange Argument) . . . 35
7.4 Exercise: Interval Covering . . . 35
7.5 Fractional Knapsack . . . 36
7.6 Interval Partitioning . . . 37

8 Huffman Coding **39**
8.1 Motivation: Variable-Length Encoding . . . 39
8.2 Prefix-Free Codes . . . 39
8.3 The Encoding Problem . . . 39
8.4 Tree Correspondence and Weighted External Path Length . . . 40
8.5 Huffman's Greedy Algorithm . . . 40
8.6 Worked Example . . . 40
8.7 Correctness Proof . . . 41
8.8 Application: Optimal Two-Way Merge . . . 42

1 Prerequisite: Asymptotic Notations

O, Ω, Θ are used to describe the asymptotic behavior of functions.

Definitions of Asymptotic Notations

(a) **Big-O (Upper Bound):** $f(n) = O(g(n))$ means that there exist positive constants c and n_0 such that $f(n) \leq cg(n)$ for all $n \geq n_0$.

$$\lim_{n\to\infty} \frac{f(n)}{g(n)} < \infty$$

(b) **Big-Omega (Lower Bound):** $f(n) = \Omega(g(n))$ means that there exist positive constants c and n_0 such that $f(n) \geq cg(n)$ for all $n \geq n_0$.

$$\lim_{n\to\infty} \frac{f(n)}{g(n)} > 0$$

(c) **Big-Theta (Tight Bound):** $f(n) = \Theta(g(n))$ means that there exist positive constants c_1, c_2 and n_0 such that $c_1 g(n) \leq f(n) \leq c_2 g(n)$ for all $n \geq n_0$.

$$0 < \lim_{n\to\infty} \frac{f(n)}{g(n)} < \infty$$

(d) **Not Comparable:** $f(n)$ and $g(n)$ are not comparable if $\lim_{n\to\infty} \frac{f(n)}{g(n)}$ does not exist.

Asymptotic Growth Hierarchy (Memorize This Order)

constant $<$ logarithmic $<$ polynomial $<$ exponential

$$O(1) \subset O(\log\log n) \subset O(\log n) \subset O(\sqrt{n}) \subset O(n) \subset O(n\log n)$$
$$\subset O(n^2) \subset O(n^k) \subset O(2^n) \subset O(n!) \subset O(n^n)$$

Common Mistakes with Asymptotic Notation

- $O(g(n))$ describes an **upper bound**, not necessarily tight. $n = O(n^2)$ is technically true but not useful.
- $\Omega(g(n))$ describes a **lower bound**. Best-case $\neq$ lower bound!
- $\Theta(g(n))$ requires **both** upper and lower bounds to match.
- Always drop constant factors and lower-order terms: $3n^2 + 5n + 7 = \Theta(n^2)$, not $\Theta(3n^2)$.

2 Divide and Conquer

Solve a problem of size n by breaking it into one or more smaller problems of size less than n. Solve the smaller problems recursively and combine their solutions to solve the large problem.

Three Steps of D&C

1. **Divide:** Split the problem into subproblems.
2. **Conquer:** Solve subproblems recursively (or directly if small enough).
3. **Combine:** Merge subproblem solutions into the original solution.

2.1 Binary Search

Given a sorted array A and a target value x, find the index of x (or -1 if not found).

Algorithm 1 Binary Search

```
function BinarySearch(A, x)
    low ← 0
    high ← length(A) − 1
    while low ≤ high do
        mid ← ⌊(low + high)/2⌋
        if A[mid] = x then
            return mid
        else if A[mid] < x then
            low ← mid + 1
        else
            high ← mid − 1
        end if
    end while
    return −1
end function
```

Binary Search Analysis

Let $T(n)$ be the number of comparisons on an array of size n.

- Recurrence: $T(n) \leq T\left(\frac{n}{2}\right) + 2$, with $T(1) = T(0) = 1$.
- Expand: $T(n) \leq T\left(\frac{n}{2}\right) + 2 \leq T\left(\frac{n}{4}\right) + 4 \leq \cdots \leq T(1) + 2\log_2 n = 1 + 2\log_2 n$.
- Therefore: $T(n) = O(\log n)$.

Algorithm 2 Binary Search (Recursive Version)

```
function BinarySearch(A, l, r, x)
    if l > r then
        return nil                                   ▷ x not present
    end if
    m ← ⌊(l + r)/2⌋
    if A[m] = x then
        return m
    else if x < A[m] then
        return BinarySearch(A, l, m − 1, x)
    else
        return BinarySearch(A, m + 1, r, x)
    end if
end function
```

Worked example: search for $x = 42$ in

$$A = [4, 7, 10, 15, 19, 20, 42, 54, 87, 90].$$

Call	Range	m	$A[m]$ vs. 42
1	$[1, 10]$	5	$19 < 42 \Rightarrow$ go right
2	$[6, 10]$	8	$54 > 42 \Rightarrow$ go left
3	$[6, 7]$	6	$20 < 42 \Rightarrow$ go right
4	$[7, 7]$	7	$A[7] = 42$ **found!**

Recursion-Tree View and the $\Omega(\log n)$ Lower Bound

The recursion tree is a single path: exactly one subproblem per level, costing 2 comparisons per internal level and 1 at the leaf, giving $2\log_2 n + 1$ comparisons — equivalent to expansion but more visual.

- Upper bound: each level halves the range, so at most $\lfloor \log_2 n \rfloor + 1$ levels $\Rightarrow T(n) = O(\log n)$.
- Lower bound: an element sitting at a deepest leaf (or a failed search) forces traversing the full path of length $\log_2 n$, so $T(n) = \Omega(\log n)$.

Hence $T(n) = \Theta(\log n)$. More generally, the decision-tree model proves no search algorithm on a sorted array can beat $\Omega(\log n)$ comparisons (there are $n + 1$ possible answers: n positions plus "not found").

2.1.1 Variant: Find the Boundary (lower_bound)

Consider the variant with **no equality test**:

Algorithm 3 Binary Search without Equality Test

```
function FindBoundary(A, l, r, x)
    if l = r then
        return l
    end if
    m ← ⌊(l + r)/2⌋
    if A[m] < x then
        return FindBoundary(A, m + 1, r, x)
    else
        return FindBoundary(A, l, m, x)
    end if
end function
```

What Does This Variant Return?

- It returns the index of the **first** element $\geq x$ — the insertion position even if x is absent (C++ `std::lower_bound`, Python `bisect.bisect_left`).
- Changing $A[m] < x$ to $A[m] \leq x$ returns the first element **strictly greater** than x (`upper_bound`, `bisect_right`); the two answers differ exactly when x occurs in A.
- Running time is still $O(\log n)$: every call halves a non-empty range.

Tail recursion into a loop. Because the recursive call is the very last action (tail recursion), it rewrites directly as a `while` loop:

```
l = 1; r = n;
while (l < r) {
    m = (l + r) / 2;
    if (A[m] < x) l = m + 1;
    else          r = m;
}
return l;
```

The iterative form avoids function-call overhead, uses $O(1)$ instead of $O(\log n)$ space, and cannot overflow the call stack — this is why library routines are written this way.

When to Use Binary Search

Binary search works whenever you can discard half the search space with one comparison. The array need not be literally sorted—it just needs a "monotonic" property that lets you decide which half contains the answer.

2.1.2 Exercise 1: Rotated Sorted Array

Let $A[1..n]$ be a sorted array of n distinct numbers that has been rotated $n-k$ steps for some unknown integer $k \in [1, n-1]$. That is, $A[1..k]$ is sorted increasing, $A[k+1..n]$ is sorted increasing, and $A[n] < A[1]$.

Example: $n = 16$, $k = 10$:

$$A = [9, 13, 16, 18, 19, 23, 28, 31, 37, 42, 0, 1, 2, 5, 7, 8]$$

(1) Find the rotation index k in $O(\log n)$. ($A[k]$ is the maximum element.)

Algorithm 4 Find Rotated Index

function FindRotatedIndex(A, l, r)
  $mid \leftarrow \left\lfloor \frac{l+r}{2} \right\rfloor$
  if $A[mid] > A[mid+1]$ **then**
    return mid
  else if $A[mid] \geq A[l]$ **then**
    return FindRotatedIndex($A, mid+1, r$)
  else
    return FindRotatedIndex($A, l, mid-1$)
  end if
end function

(2) Search for x in $O(\log n)$.

Algorithm 5 Search in Rotated Sorted Array

```
function SearchRotated(A, x, n)
    k ← FindRotatedIndex(A, 1, n)
    if x ≥ A[1] then
        return BinarySearch(A[1..k], x)
    else
        return BinarySearch(A[k + 1..n], x)
    end if
end function
```

2.1.3 Exercise 2: Finding the Last 0

Array $A[1..n]$ contains 0s followed by 1s.

(1) Standard $O(\log n)$ solution.

Algorithm 6 Find Last 0

```
function FindLast0(A, l, r)
    mid ← ⌊(l + r)/2⌋
    if A[mid] = 0 and A[mid + 1] = 1 then
        return mid
    else if A[mid] = 0 then
        return FindLast0(A, mid + 1, r)
    else
        return FindLast0(A, l, mid − 1)
    end if
end function
```

(2) $O(\log k)$ solution when $k \ll n$. Exponential search (gallop): double the index until we overshoot, then binary search.

Algorithm 7 Find Last 0 with $O(\log k)$

```
i ← 1
while i < n ∧ A[i] = 0 do
    i ← min{2i, n}
end while
if A[i] = 0 then
    return i                          ▷ all zeros: last 0 is n
else if i = 1 then
    return −1                         ▷ no zero
else
    return FindLast0(A, ⌊i/2⌋, i)
end if
```

2.2 Mergesort & Inversion Count

2.2.1 Warm-up: Tower of Hanoi

Three pegs, n disks of decreasing size. Move all disks from source to target peg.

Hanoi Recurrence

Let $T(n)$ = minimum moves for n disks.

- Basis: $T(1) = 1$.
- Recursive step: move $n - 1$ disks source→aux, move largest source→target, move

$n-1$ aux→target.

$$T(n) = 2T(n-1) + 1$$

- Solution: $T(n) = 2^n - 1$. (Expand: $T(n) = 2^k T(n-k) + (2^k - 1)$, set $k = n-1$.)

Algorithm 8 Hanoi Tower

```
function HANOI(n, A, B, C)
    if n = 1 then
        Move disk from A to C
    else
        HANOI(n − 1, A, C, B)
        Move disk from A to C
        HANOI(n − 1, B, A, C)
    end if
end function
```

2.2.2 Merge: Combining Two Sorted Halves

Given two sorted subarrays $A[p..q]$ and $A[q+1..r]$, merge them into $A[p..r]$ sorted.

Algorithm 9 Merge

```
function MERGE(A, p, q, r)
    n1 ← q − p + 1;    n2 ← r − q
    Create arrays L[1..n1 + 1], R[1..n2 + 1]
    for i = 1 to n1 do
        L[i] ← A[p + i − 1]
    end for
    for j = 1 to n2 do
        R[j] ← A[q + j]
    end for
    L[n1 + 1] ← ∞;    R[n2 + 1] ← ∞                ▷ sentinels
    i ← 1;    j ← 1
    for k = p to r do
        if L[i] ≤ R[j] then
            A[k] ← L[i];    i ← i + 1
        else
            A[k] ← R[j];    j ← j + 1
        end if
    end for
end function
```

Merge Time Complexity

Merge runs in $\Theta(n)$ time on $n = r-p+1$ elements: two copy loops ($n_1+n_2 = n$) plus one comparison loop (n iterations). Sentinels avoid boundary checks — clean but not cache-friendly in practice.

Worked example: merge the sorted halves $L = [2, 4, 5, 7, \infty]$ and $R = [1, 2, 3, 6, \infty]$:

k	compare	winner	$A[p \dots r]$ so far
1	$L[1] = 2$ vs $R[1] = 1$	$R[1]$	[1]
2	$L[1] = 2$ vs $R[2] = 2$	$L[1]$ ($\leq$ keeps it stable)	[1, 2]
3	$L[2] = 4$ vs $R[2] = 2$	$R[2]$	[1, 2, 2]
4	$L[2] = 4$ vs $R[3] = 3$	$R[3]$	[1, 2, 2, 3]
5	$L[2] = 4$ vs $R[4] = 6$	$L[2]$	[1, 2, 2, 3, 4]
6	$L[3] = 5$ vs $R[4] = 6$	$L[3]$	[1, 2, 2, 3, 4, 5]
7	$L[4] = 7$ vs $R[4] = 6$	$R[4]$	[1, 2, 2, 3, 4, 5, 6]
8	$L[4] = 7$ vs ∞	$L[4]$	[1, 2, 2, 3, 4, 5, 6, 7]

Each step advances the head pointer (i or j) of the list whose element was copied; exactly one element is placed per iteration.

2.2.3 Mergesort Algorithm

Algorithm 10 Mergesort

function MERGESORT(A, p, r)
  if $p < r$ **then**
    $q \leftarrow \left\lfloor \frac{p+r}{2} \right\rfloor$ ▷ Divide
    MERGESORT(A, p, q) ▷ Conquer left
    MERGESORT($A, q+1, r$) ▷ Conquer right
    MERGE(A, p, q, r) ▷ Combine
  end if
end function

Mergesort Analysis

Recurrence for n elements (power of 2 for simplicity):

$$T(n) = \begin{cases} 1 & \text{if } n = 1, \\ 2T\left(\frac{n}{2}\right) + \Theta(n) & \text{if } n > 1. \end{cases}$$

Solution: $T(n) = \Theta(n \log n)$. (Master Theorem: $a = 2, b = 2, f(n) = \Theta(n)$, case 2.)

Worked example: $A = [5, 2, 4, 7, 1, 3, 2, 6]$.

Stage	Array state
input	$[5, 2, 4, 7 \mid 1, 3, 2, 6]$
after sorting left half recursively	$[2, 4, 5, 7 \mid 1, 3, 2, 6]$
after sorting right half recursively	$[2, 4, 5, 7 \mid 1, 2, 3, 6]$
after $\Theta(n)$ merge	$[1, 2, 2, 3, 4, 5, 6, 7]$

Solving the Recurrence by Expansion / Recursion Tree

Expanding directly:

$$\begin{aligned} T(n) &= 2T(n/2) + n = 2^2 T(n/4) + 2n = 2^3 T(n/8) + 3n \\ &= 2^i T(n/2^i) + in = nT(1) + n \log_2 n = n \log n + n. \end{aligned}$$

On the recursion tree, level i has 2^i nodes, each paying $n/2^i$, so the cost of every level is $2^i \cdot (n/2^i) = n$; with $\log_2 n + 1$ levels the total is $n(\log_2 n + 1)$.

Why $\Omega(n \log n)$ Too?

The merge step costs $\Theta(n)$ on every input — there is no lucky case that skips work. Replacing each O by Ω in the analysis is still valid, so $T(n) = \Omega(n \log n)$ and overall $T(n) = \Theta(n \log n)$: every input is a worst-case input for mergesort.

Mergesort Properties

- **Stable:** yes (preserves order of equal keys if $L[i] \leq R[j]$ uses $\leq$, not $<$).
- **In-place:** no (needs auxiliary arrays L and R; $\Theta(n)$ extra space).
- **Best / Worst / Average:** all $\Theta(n \log n)$ (unlike Quicksort).

2.2.4 Counting Inversions

An **inversion** is a pair (i, j) with $i < j$ and $A[i] > A[j]$. Count the total number of inversions.

Brute Force is Too Slow

Checking all $\binom{n}{2}$ pairs takes $\Theta(n^2)$ time. With D&C we can do $\Theta(n \log n)$ by **piggybacking on Mergesort**.

Algorithm 11 Count Inversions via Mergesort

```
function CountInv(A, p, r)
    if p ≥ r then
        return 0
    end if
    q ← ⌊(p + r)/2⌋
    left ← CountInv(A, p, q)
    right ← CountInv(A, q + 1, r)
    cross ← MergeCount(A, p, q, r)
    return left + right + cross
end function

function MergeCount(A, p, q, r)
    ... same as Merge, but whenever R[j] is copied before L[i] is exhausted:
    add (n_1 − i + 1) to cross (all remaining L[i..n_1] are > R[j])
    return cross
end function
```

Why Cross-Inversions Work

When copying $R[j]$ before $L[i]$ is empty, **every remaining element in L forms an inversion with $R[j]$** because L indices $< R$ indices (by array split) and $L[i..n_1] > R[j]$ (since we picked $R[j]$ first). This adds exactly $n_1 - i + 1$ inversions in $O(1)$.

Applications. The inversion number measures (i) how "sorted" an array is (0 = sorted, $\binom{n}{2}$ = reverse sorted), and (ii) the similarity between two rankings: map one ranking to the identity permutation and count inversions in the other.

Worked example: $A = [4, 8, 10, 2, 1, 5, 12, 11, 3, 7, 6, 9]$.

Split into blue half $[4, 8, 10, 2, 1, 5]$ and green half $[12, 11, 3, 7, 6, 9]$:

- 5 blue-blue inversions and 8 green-green inversions (found recursively).
- 9 blue-green inversions: $(5, 3)$, $(4, 3)$, $(8, 6)$, $(8, 3)$, $(8, 7)$, $(10, 6)$, $(10, 9)$, $(10, 3)$, $(10, 7)$.
- Total: $5 + 8 + 9 = \mathbf{22}$.

Counting cross inversions during the merge. With 1-indexed arrays split at m, when $R[j]$ is copied while the L-pointer is at i, the blue elements already passed ($L[1..i-1]$) are $\le R[j]$; every remaining blue element $L[i..m]$ is $> R[j]$, so the contribution of $R[j]$ is

$$I[j] = m - i + 1 \quad \big(\text{or, with the slide's indexing, } m - l - i + 2\big).$$

Tiny trace: sorted halves $L = [1, 3, 5, 7]$, $R = [2, 4, 6, 8]$:

Element copied from R	pointer i	$L[i..]$ greater than it	$I[j]$
2	1 ($L[1] = 1 \le 2$ passed, then $i = 2$)	3, 5, 7	3
4	3	5, 7	2
6	4	7	1
8	past end	—	0

Cross inversions $= 3+2+1+0 = 6$; the merge simultaneously produces $[1, 2, 3, 4, 5, 6, 7, 8]$, preserving the sortedness invariant needed by the parent call.

Recursion-tree accounting. Each leaf returns 0; each internal node returns its cross count from the merge, and a node reports the sum of its two children plus its own cross count. For the 12-element example the two size-6 root children return 5 (blue half) and 8 (green half), and the root merge contributes 9:

root: $5 + 8 + 9 = 22$
left child `Sort-and-Count`$[4, 8, 10, 2, 1, 5] \to 5$
right child `Sort-and-Count`$[12, 11, 3, 7, 6, 9] \to 8$
(their values are sums of cross counts at lower levels; leaves contribute 0)

Since both the counting and the merging are linear scans, `Merge-and-Count` costs $\Theta(n)$ and the overall recurrence is again $T(n) = 2T(n/2) + \Theta(n) = \Theta(n \log n)$.

Algorithm 12 Merge-and-Count (precondition: $A[l..m]$ and $A[m+1..r]$ sorted)

copy $A[l..m]$ into $L[1..n_1]$, $A[m+1..r]$ into $R[1..n_2]$, append ∞ sentinels
$i \leftarrow 1$; $j \leftarrow 1$; $c \leftarrow 0$
for $k = l$ **to** r **do**
  if $L[i] \le R[j]$ **then**
    $A[k] \leftarrow L[i]$; $i \leftarrow i + 1$
  else
    $A[k] \leftarrow R[j]$; $j \leftarrow j + 1$
    $c \leftarrow c + (n_1 - i + 1)$ ▷ all remaining L elements invert $R[j]$
  end if
end for
return c

2.3 Maximum Subarray

Given an array $A[1..n]$, find indices i, j maximizing $\sum_{t=i}^{j} A[t]$ (the empty subarray is not allowed). (Useful for best-buy/best-sell stock problem.)

Motivating example. Yearly profits (M$) $[3, 2, 1, -7, 5, 2, -1, 3, -1]$: the best span is years 5–8 with $5 + 2 - 1 + 3 = 9$ M$.

Two Simpler Solutions First

1. **Brute force $\Theta(n^3)$:** enumerate all pairs (i, j) and re-sum $A[i..j]$ with an inner loop. The j- and k-loops average $n/2$ and $n/3$ iterations.
2. **Data reuse $\Theta(n^2)$:** exploit $V(i, j) = V(i, j-1) + A[j]$ — extend the running sum instead of restarting, removing the inner loop.

The D&C solution improves this to $\Theta(n \log n)$, and Kadane's DP reaches $\Theta(n)$.

D&C Strategy

For a subarray $A[low..high]$ with midpoint mid, the maximum subarray is **exactly one** of:
1. Entirely in $A[low..mid]$ (recurse left).
2. Entirely in $A[mid + 1..high]$ (recurse right).
3. **Crossing the midpoint:** one end in left half, one end in right half (compute directly).

Return the maximum of the three.

Algorithm 13 Maximum Subarray (D&C)

```
function MaxSubarray(A, low, high)
    if low = high then
        return (low, high, A[low])
    end if
    mid ← ⌊(low + high)/2⌋
    (left_lo, left_hi, left_sum) ← MaxSubarray(A, low, mid)
    (right_lo, right_hi, right_sum) ← MaxSubarray(A, mid + 1, high)
    (cross_lo, cross_hi, cross_sum) ← MaxCrossing(A, low, mid, high)
    if left_sum ≥ right_sum and left_sum ≥ cross_sum then
        return (left_lo, left_hi, left_sum)
    else if right_sum ≥ left_sum and right_sum ≥ cross_sum then
        return (right_lo, right_hi, right_sum)
    else
        return (cross_lo, cross_hi, cross_sum)
    end if
end function

function MaxCrossing(A, low, mid, high)
    left_sum ← −∞;    sum ← 0
    for i = mid downto low do
        sum ← sum + A[i]
        if sum > left_sum then
            left_sum ← sum;    max_left ← i
        end if
    end for
    right_sum ← −∞;    sum ← 0
    for j = mid + 1 to high do
        sum ← sum + A[j]
        if sum > right_sum then
            right_sum ← sum;    max_right ← j
        end if
    end for
    return (max_left, max_right, left_sum + right_sum)
end function
```

Max Subarray Analysis

Recurrence:

$$T(n) = \begin{cases} \Theta(1) & n = 1, \\ 2T(n/2) + \Theta(n) & n > 1. \end{cases}$$

Same as Mergesort: $T(n) = \Theta(n \log n)$.

Kadane's Algorithm (iterative DP, not D&C): $O(n)$ time, $O(1)$ space.

$$\text{ending_here}_i = \max\big(A[i],\ \text{ending_here}_{i-1} + A[i]\big),$$
$$\text{best}_i = \max\big(\text{best}_{i-1},\ \text{ending_here}_i\big).$$

Kadane trace on $A = [-2, 1, -3, 4, -1, 2, 1, -5, 4]$ (start $V = V_{\max} = -\infty$):

i	$A[i]$	$V = \max(V + A[i], A[i])$ (best ending at i)	$V_{\max}$	best window so far
1	-2	-2	-2	$[-2]$
2	1	1 (restart)	1	$[1]$
3	-3	-2	1	$[1]$
4	4	4 (restart)	4	$[4]$
5	-1	3	4	$[4]$
6	2	5	5	$[4, -1, 2]$
7	1	6	**6**	$[4, -1, 2, 1]$
8	-5	1	6	$[4, -1, 2, 1]$
9	4	5	6	$[4, -1, 2, 1]$

Answer: 6, achieved by $[4, -1, 2, 1]$. The "restart or extend" choice at every step is the DP decision.

2.3.1 Application: Maximizing Stock Profit

Given prices $p[1..n]$, buy once and sell later ($i < j$) to maximize $p[j] - p[i]$ (a trade is mandatory).

Direct D&C: split at m. The optimal pair is either (1) entirely in the left half, (2) entirely in the right half, or (3) **crossing**: buy at the minimum price in $p[l..m]$ and sell at the maximum price in $p[m+1..r]$, found by two linear scans. Same recurrence $2T(n/2) + \Theta(n) = \Theta(n \log n)$.

Reduction to Maximum Subarray ($O(n)$ preprocessing)

Build the difference array of consecutive-day changes:

$$\text{Profit}[i] = p[i+1] - p[i].$$

Every interval telescopes:

$$\sum_{t=i}^{j-1} \text{Profit}[t] = \big(p[i+1] - p[i]\big) + \big(p[i+2] - p[i+1]\big) + \cdots + \big(p[j] - p[j-1]\big) = p[j] - p[i].$$

So the maximum-profit trade is exactly the maximum subarray of `Profit`, solvable by Kadane in $\Theta(n)$.

Example: prices $[3, 6, 10, 1, 8, 2, 6, 5, 10, 3, 1, 9]$ give differences

$$[3, 4, -9, 7, -6, 4, -1, 5, -7, -2, 8],$$

whose maximum subarray identifies the best buy/sell pair.

2.4 Multiplication and Master Theorem

2.4.1 Integer Multiplication (Karatsuba)

Given two n-bit integers x and y, compute $x \cdot y$.

Grade-School Multiplication

The standard pencil-and-paper algorithm is $\Theta(n^2)$: multiply each bit of y by x (shift & add), giving n rows of $O(n)$ bits each. Example: 163×97:

$$\begin{array}{rrr} & & 163 \\ \times & & 97 \\ \hline & 1141 & (163 \times 7) \\ + \quad 1467 & & (163 \times 9, \text{ shifted}) \\ \hline & & 15811 \end{array}$$

In binary, 10100011×01100001 similarly produces n shifted partial products: $\Theta(n^2)$ single-bit multiplications and $\Theta(n^2)$ binary additions.

2.4.2 First D&C Attempt: Four Half-Size Products

Split $a = 2^{n/2}a_1 + a_0$, $b = 2^{n/2}b_1 + b_0$ (high/low $n/2$ bits):

$$a \cdot b = 2^n a_1 b_1 + 2^{n/2}(a_1 b_0 + a_0 b_1) + a_0 b_0.$$

Multiplication by 2^k is just a k-bit left shift (e.g. $12 = 00001100$, then $12 \cdot 2^3 = 96 = 01100000$), so the four products $a_1b_1, a_1b_0, a_0b_1, a_0b_0$ are computed recursively and combined with $\Theta(n)$ additions:

$$T(n) = 4T(n/2) + \Theta(n), \qquad T(1) = 1.$$

Analysis: The Naive D&C Is Still $\Theta(n^2)$

Expanding: $T(n) = 4^i T(n/2^i) + n(2^{i-1} + \cdots + 2 + 1)$; at $i = \log_2 n$,

$$T(n) = n^2 + n(2^{\log_2 n} - 1) = n^2 + n(n-1) = \Theta(n^2).$$

On the recursion tree, level i costs $4^i \cdot n/2^i = n2^i$, so costs double each level and the leaves dominate ($4^{\log_2 n} = n^2$). Essentially the method still multiplies every bit of a by every bit of b. The contrast with mergesort: one problem spawns **4** subproblems of size $n/2$, not 2.

Karatsuba Split

Write $x = a \cdot 2^{n/2} + b$, $y = c \cdot 2^{n/2} + d$ (split into high/low halves). Then:

$$x \cdot y = ac \cdot 2^n \; + \; \big[(a+b)(c+d) - ac - bd\big] \cdot 2^{n/2} \; + \; bd.$$

Instead of 4 half-size multiplications (ac, ad, bc, bd), we use **3**: ac, bd, and $(a+b)(c+d)$, subtracting to recover $ad + bc$.

Algorithm 14 Karatsuba Multiplication

```
function MULTIPLY(x, y, n)
    if n = 1 then
        return x · y
    end if
    a ← high n/2 bits of x;    b ← low n/2 bits of x
    c ← high n/2 bits of y;    d ← low n/2 bits of y
    ac ← MULTIPLY(a, c, n/2)
    bd ← MULTIPLY(b, d, n/2)
    ab_cd ← MULTIPLY(a + b, c + d, n/2)
    return ac ≪ n + (ab_cd − ac − bd) ≪ (n/2) + bd
end function
```

Karatsuba Recurrence

$$T(n) = 3T\left(\frac{n}{2}\right) + \Theta(n).$$

By Master Theorem case 1 ($\log_b a = \log_2 3 \approx 1.585 > 1$):

$$T(n) = \Theta(n^{\log_2 3}) \approx \Theta(n^{1.585}),$$

asymptotically better than $\Theta(n^2)$.

Decimal illustration ($n = 2$ digits, no recursion needed). Multiply 23×14: $a_1 = 2, a_0 = 3$, $c_1 = 1, d_0 = 4$.

$$ac = 2 \cdot 1 = 2, \qquad bd = 3 \cdot 4 = 12,$$
$$(a+b)(c+d) = 5 \cdot 5 = 25, \qquad ad + bc = 25 - 2 - 12 = 11,$$
$$23 \cdot 14 = 2 \cdot 100 + 11 \cdot 10 + 12 = 200 + 110 + 12 = \mathbf{322}.$$

Only **3** small multiplications $(2, 12, 25)$ instead of 4.

History and Practice

Karatsuba (1960) was the first algorithm to beat $\Theta(n^2)$. Splitting into r parts and using fewer than r^2 products gives the Toom–Cook family: Toom-3 $\Rightarrow \Theta(n^{\log_3 5}) = \Theta(n^{1.465})$. FFT-based methods reached $\Theta(n \log n \log\log n)$; a $\Theta(n \log n)$ integer-multiplication algorithm was proven by Harvey & van der Hoeven in 2019. Real big-integer libraries combine methods by input size: textbook for tiny inputs, Karatsuba/Toom for medium, FFT for very large.

2.4.3 Master Theorem for Recurrences

For recurrences of the form:

$$T(n) = a\,T\left(\frac{n}{b}\right) + f(n), \quad a \geq 1,\ b > 1.$$

Compare $f(n)$ with $n^{\log_b a}$ (the "leaf work").

Master Theorem — Three Cases

1. **Case 1 (Leaves dominate):** If $f(n) = O(n^{\log_b a - \epsilon})$ for some $\epsilon > 0$,
$$T(n) = \Theta(n^{\log_b a}).$$
2. **Case 2 (Balanced):** If $f(n) = \Theta(n^{\log_b a} \log^k n)$ for some $k \geq 0$,
$$T(n) = \Theta(n^{\log_b a} \log^{k+1} n).$$
3. **Case 3 (Root dominates):** If $f(n) = \Omega(n^{\log_b a + \epsilon})$ and the **regularity condition** $a\,f(n/b) \leq c\,f(n)$ holds for some $c < 1$ and large n,
$$T(n) = \Theta(f(n)).$$

Common D&C Recurrences (Memorize!)

Algorithm	Recurrence	Solution
Binary search	$T(n) = T(n/2) + \Theta(1)$	$\Theta(\log n)$ (Case 2, $k = 0$)
Mergesort	$T(n) = 2T(n/2) + \Theta(n)$	$\Theta(n \log n)$ (Case 2, $k = 0$)
Karatsuba	$T(n) = 3T(n/2) + \Theta(n)$	$\Theta(n^{\log_2 3})$ (Case 1)
Strassen matrix	$T(n) = 7T(n/2) + \Theta(n^2)$	$\Theta(n^{\log_2 7})$ (Case 1)
Bad D&C (naive 4-mult)	$T(n) = 4T(n/2) + \Theta(n)$	$\Theta(n^2)$ (Case 1)

Why the Master Theorem Works (Recursion-Tree Intuition)

For the simplified form $T(n) = aT(n/b)+n$: level i contains a^i subproblems of size n/b^i and costs $n(a/b)^i$; there are $\log_b n$ levels and the leaves alone contribute $a^{\log_b n} = n^{\log_b a}$. Total:

$$T(n) = n^{\log_b a} + n \sum_{i=0}^{\log_b n - 1} \left(\frac{a}{b}\right)^i.$$

The geometric ratio a/b decides everything:

- $a > b$: increasing geometric series, **leaves dominate**, $T(n) = \Theta(n^{\log_b a})$.
- $a = b$: every term is 1, each level costs n, $T(n) = \Theta(n \log n)$.
- $a < b$: decreasing geometric series is $\Theta(1)$, **the root dominates**, $T(n) = \Theta(n)$.

2.5 Matrix Multiplication and Strassen's Algorithm

Given two $n \times n$ matrices, compute $C = AB$ with $c_{ij} = \sum_{k=1}^{n} a_{ik} b_{kj}$. The brute-force triple loop costs $\Theta(n^3)$ (n^2 entries, each a length-n dot product).

Naive D&C: partition each matrix into four $n/2 \times n/2$ blocks:

$$\begin{pmatrix} C_{11} & C_{12} \\ C_{21} & C_{22} \end{pmatrix} = \begin{pmatrix} A_{11} & A_{12} \\ A_{21} & A_{22} \end{pmatrix} \begin{pmatrix} B_{11} & B_{12} \\ B_{21} & B_{22} \end{pmatrix},$$

requiring **8** block multiplications (e.g. $C_{11} = A_{11}B_{11} + A_{12}B_{21}$) and $\Theta(n^2)$ block additions:

$$T(n) = 8T(n/2) + \Theta(n^2) = \Theta(n^3)$$

(Master case 1: $\log_2 8 = 3$, and $n^{\log_2 8} = n^3$ dominates n^2). No improvement.

Strassen's Idea: 7 Block Products

Strassen (1969) computes the 2×2 block product with only **7** recursive multiplications plus $\Theta(n^2)$ additions/subtractions:

$$\begin{aligned} P_1 &= A_{11}(B_{12} - B_{22}), & P_2 &= (A_{11} + A_{12})B_{22}, & P_3 &= (A_{21} + A_{22})B_{11}, \\ P_4 &= A_{22}(B_{21} - B_{11}), & P_5 &= (A_{11} + A_{22})(B_{11} + B_{22}), & & \\ P_6 &= (A_{12} - A_{22})(B_{21} + B_{22}), & P_7 &= (A_{11} - A_{21})(B_{11} + B_{12}), & & \end{aligned}$$

then

$$\begin{aligned} C_{11} &= P_5 + P_4 - P_2 + P_6, & C_{12} &= P_1 + P_2, \\ C_{21} &= P_3 + P_4, & C_{22} &= P_5 + P_1 - P_3 - P_7. \end{aligned}$$

Hence

$$T(n) = 7T(n/2) + \Theta(n^2) = \Theta(n^{\log_2 7}) = \Theta(n^{2.807\ldots}),$$

the first algorithm to beat cubic matrix multiplication. It is practical only for large dense matrices ($n \gtrsim 100$): the extra additions and recursion overhead make it slower on small matrices.

3 Randomized Algorithms

3.1 Why Randomization?

A randomized algorithm flips coins during execution. Worst-case analysis becomes **expected**-case analysis, often beating deterministic worst-case bounds.

Two Types of Randomized Algorithms

- **Las Vegas:** always correct; running time is a random variable (e.g. randomized Quicksort). We bound the expected runtime.
- **Monte Carlo:** always fast; has a small probability of being wrong (e.g. randomized min-cut / Karger). Repeat to reduce error.

How We Measure Randomized Running Time

We take the **expected running time on the worst-case input**: $\max_{I:\,|I|=n} \mathbb{E}_R[T(I,R)]$, where the expectation is only over the algorithm's internal randomness R (random pivots, random permutations, ...). Randomized algorithms are often simpler and faster in practice; whether randomization can give an asymptotic advantage over deterministic algorithms remains open for some problems (e.g. $O(1)$-time hash tables with $O(n)$ space).

3.2 Probability Review

For a discrete random variable X taking values with probabilities $\Pr[X = i]$:

$$\mathbb{E}[X] = \sum_i i \cdot \Pr[X = i].$$

Dice warm-ups:

- One fair 6-sided die: $\mathbb{E}[X] = \frac{1+2+\cdots+6}{6} = 3.5$.
- Sum of two dice: by linearity of expectation, $\mathbb{E}[X_1 + X_2] = 3.5 + 3.5 = \mathbf{7}$.
- Maximum of two dice: $\Pr[\max = k] = (k^2 - (k-1)^2)/36 = (2k-1)/36$, so

$$\mathbb{E}[\max] = 1 \cdot \tfrac{1}{36} + 2 \cdot \tfrac{3}{36} + 3 \cdot \tfrac{5}{36} + \cdots + 6 \cdot \tfrac{11}{36} \approx 4.47.$$

Two useful facts: $\mathbb{E}[XY] = \mathbb{E}[X]\mathbb{E}[Y]$ only when X, Y are **independent** (but linearity of expectation needs no independence — that is why it is so powerful). For a non-negative integer-valued variable:

$$\mathbb{E}[X] = \sum_{i \geq 1} \Pr[X \geq i]$$

(the "tail-sum" formula, handy for waiting times).

3.3 Linearity of Expectation

For any random variables $X_1, X_2, \ldots, X_n$ (they need **not** be independent):

$$\mathbb{E}\left[\sum_{i=1}^{n} X_i\right] = \sum_{i=1}^{n} \mathbb{E}[X_i].$$

This is the single most useful tool for analyzing randomized algorithms.

Indicator Variables

To count "how many events happened", set:

$$X_i = \mathbf{1}\{\text{event } i \text{ occurs}\} = \begin{cases} 1 & \text{if event } i \text{ happens,} \\ 0 & \text{otherwise.} \end{cases}$$

Then $\mathbb{E}[X_i] = \Pr(\text{event } i)$, and $X = \sum X_i$ gives the total count with $\mathbb{E}[X] = \sum \Pr(\text{event } i)$.

3.4 Example: Hiring Problem

Interview candidates in random order; hire whenever candidate is better than all previous. How many hires (expected)?

Let $X_i = 1$ if candidate i is hired (i.e. best among first i candidates). In random order this occurs with probability $1/i$.

$$\mathbb{E}[X] = \sum_{i=1}^{n} \mathbb{E}[X_i] = \sum_{i=1}^{n} \frac{1}{i} = H_n = \Theta(\log n).$$

So the expected number of hires is the n-th harmonic number, $\ln n + O(1)$.

3.5 More Indicator-Variable Examples

3.5.1 Guessing Shuffled Cards

Shuffle n cards and guess each card as it is revealed.

- **No memory:** each turn guess a uniformly random card from the full deck. Let $X_i = \mathbf{1}\{\text{guess } i \text{ correct}\}$; $\mathbb{E}[X_i] = 1/n$ regardless of other guesses, so $\mathbb{E}[X] = \sum_i 1/n = \mathbf{1}$.
- **With memory:** guess uniformly among the unseen cards. At turn i there are $n-i+1$ candidates, so $\mathbb{E}[X_i] = 1/(n-i+1)$ and

$$\mathbb{E}[X] = \frac{1}{n} + \frac{1}{n-1} + \cdots + \frac{1}{1} = H_n = \Theta(\log n).$$

Note that the events are dependent in both games — linearity of expectation applies anyway.

3.5.2 Hash Table with Chaining

A hash function $h : U \to \{0, \ldots, m-1\}$ maps n keys into a table of m slots; collisions are resolved by a linked list per slot. Under the **simple uniform hashing assumption** ($h(x)$ uniform and mutually independent across keys), define $I_y = \mathbf{1}\{h(y) = h(x)\}$ for each stored key $y \neq x$. Then $\mathbb{E}[I_y] = 1/m$ and

$$\mathbb{E}[\text{keys colliding with } x] = \sum_{y \neq x} \mathbb{E}[I_y] \leq \frac{n}{m}.$$

With load factor $\alpha = n/m$, insert/delete/query take $O(1+\alpha)$ expected time; choosing $m = \Theta(n)$ makes this $O(1)$. (This is the idea behind C++ `unordered_map`.)

3.5.3 Birthday Paradox

With k people and $n = 365$ equiprobable birthdays, let $X_{ij} = \mathbf{1}\{i, j \text{ share a birthday}\}$ ($\mathbb{E}[X_{ij}] = 1/n$):

$$\mathbb{E}[X] = \binom{k}{2}\frac{1}{n} = \frac{k(k-1)}{2n}.$$

Expecting ≥ 1 collision means $k(k-1)/(2n) \geq 1$, i.e. $k \gtrsim \sqrt{2n} + 1$ — about **28** people for birthdays, far below the intuitive $365/2$. (This threshold also dictates when an m-slot hash table starts seeing collisions.)

3.6 Generating a Uniform Random Permutation

The hiring analysis assumes a uniformly random candidate order; the Fisher–Yates shuffle produces one in place:

Algorithm 15 Fisher–Yates Shuffle

```
function RandomPermute(A)
    for i = 1 to n do
        swap A[i] ↔ A[Random(1, i)]
    end for
end function
```

$O(n)$ time and $O(1)$ extra space.

Correctness: Every Permutation Has Probability $1/n!$

Induct on i: after iteration i, every permutation of $A[1..i]$ occurs with probability $1/i!$. For target $(a_1, \ldots, a_i)$, let j be the position of the original $A[i]$ in this target. The target appears iff (i) after iteration $i-1$ the other $i-1$ elements had the required order (probability $1/(i-1)!$ by induction), and (ii) $\text{Random}(1, i) = j$ (probability $1/i$). Multiplying: $1/(i-1)! \cdot 1/i = 1/i!$.

3.7 Geometric Distribution and Coupon Collector

Waiting for a success. If each coin flip succeeds with probability p, the number X of flips to the first head is geometric:

$$\Pr[X = j] = (1-p)^{j-1}p, \qquad \mathbb{E}[X] = \sum_{i \geq 1}\Pr[X \geq i] = \sum_{i \geq 1}(1-p)^{i-1} = \frac{1}{p}.$$

Coupon collector. Each cereal box contains one of n coupon types uniformly at random; how many boxes until all types are seen? Split the process into stages: after collecting i distinct coupons, the chance the next box is new is $(n-i)/n$, so stage i is geometric with mean $n/(n-i)$:

$$\mathbb{E}[X] = \sum_{i=0}^{n-1}\frac{n}{n-i} = n\sum_{j=1}^{n}\frac{1}{j} = nH_n = \Theta(n \log n).$$

3.8 Case Study: The Majority Element

Find the element appearing $> n/2$ times, assuming it exists.

Method	Time	Space	Idea
Sort and count	$O(n \log n)$	$O(1)$	majority, if present, occupies the middle position
Divide and conquer	$O(n \log n)$	$O(\log n)$	winners of halves, equality tests only
Hashing	$O(n)$ expected	$O(n)$	frequency map
Try-and-check	$O(n)$ expected	$O(1)$	sample a random element, verify by counting; retry
Boyer–Moore vote	$O(n)$ worst-case	$O(1)$	one cancellation pass, equality tests only

Algorithm 16 Boyer–Moore Majority Vote

```
c ← 0
for i = 1 to n do
    if c = 0 then
        m ← A[i];   c ← 1
    else if m = A[i] then
        c ← c + 1
    else
        c ← c − 1                    ▷ cancel A[i] against one stored copy of m
    end if
end for
return m
```

Why the Vote Works

Invariant: after processing $A[1..i]$, the prefix equals "c copies of m plus elements already cancelled in pairs." Each cancellation removes two different items, so any single value can lose at most $n/2$ copies. A true majority ($> n/2$ copies) can never be completely cancelled — it must survive as m. (If a majority is not guaranteed, a second $O(n)$ verification pass is required.)

4 Quicksort

4.1 Deterministic Quicksort (Lomuto Partition)

Pick $A[r]$ as pivot; partition into $<$ pivot (left), $=$ pivot (middle), $>$ pivot (right).

Algorithm 17 Quicksort

```
function QUICKSORT(A, p, r)
    if p < r then
        q ← PARTITION(A, p, r)
        QUICKSORT(A, p, q − 1)
        QUICKSORT(A, q + 1, r)
    end if
end function

function PARTITION(A, p, r)
    pivot ← A[r]
    i ← p − 1
    for j = p to r − 1 do
        if A[j] ≤ pivot then
            i ← i + 1
            swap A[i] ↔ A[j]
        end if
    end for
    swap A[i + 1] ↔ A[r]
    return i + 1
end function
```

Partition Invariants

At the end of each j-iteration in Partition:

- $A[p \dots i] \le pivot$,
- $A[i+1 \dots j] > pivot$,
- $A[j+1 \dots r-1]$ are unprocessed,
- $A[r] = pivot$.

Maintain with swaps; runtime is $\Theta(r - p + 1)$ on subarray size n.

Worked example: Partition $[2, 8, 7, 1, 3, 5, 6, 4]$ with pivot $A[r] = 4$ (Lomuto; bold marks $A[i]$ boundary):

j	$A[j] \le 4$?	i after step	array
start	—	$p-1$	$[2, 8, 7, 1, 3, 5, 6, \mathbf{4}]$
p ($A[j]=2$)	yes	p	$[\mathbf{2}, 8, 7, 1, 3, 5, 6, 4]$
$p+1$ (8)	no	p	$[2, 8, 7, 1, 3, 5, 6, 4]$
$p+2$ (7)	no	p	$[2, 8, 7, 1, 3, 5, 6, 4]$
$p+3$ (1)	yes	$p+1$	$[2, \mathbf{1}, 7, 8, 3, 5, 6, 4]$
$p+4$ (3)	yes	$p+2$	$[2, 1, \mathbf{3}, 8, 7, 5, 6, 4]$
$p+5, p+6$ (5, 6)	no	$p+2$	unchanged
final swap	—	$q = p+3$	$[2, 1, 3, \mathbf{4}, 7, 5, 6, 8]$

Pivot lands at its final sorted position; quicksort then recurses on $[2, 1, 3]$ and $[7, 5, 6, 8]$ independently.

4.2 Quicksort Analysis

Deterministic Quicksort Worst-Case

If the input is **already sorted** and pivot is always the **last** (or first) element:

$$T(n) = T(n-1) + \Theta(n) \quad \Longrightarrow \quad T(n) = \Theta(n^2).$$

Same as insertion sort! This is why we randomize.

Best case: perfectly balanced splits give $T(n) = 2T(n/2) + \Theta(n) = \Theta(n \log n)$.

"Balanced Enough" Splits: 1/4–3/4

Suppose every split were a **worst balanced case**, sizes $n/4$ and $3n/4$:

$$T(n) = T(n/4) + T(3n/4) + cn.$$

The recursion tree has a short left spine (depth $\log_4 n$) and a long right spine (depth $\log_{4/3} n$); every level does at most cn work, so even here

$$T(n) \le cn \cdot \log_{4/3} n = O(n \log n).$$

With a random pivot, Pr[pivot rank lies in middle half $[n/4, 3n/4]$] $= 1/2$ — so on average every other split is good, which is already enough for $\Theta(n \log n)$ expected depth.

Randomized Quicksort Expected Time

Pick pivot **uniformly at random** from $A[p..r]$ (swap it to position r before Partition). Let X = total number of **comparisons** made. Use indicator variables $X_{ij} = 1$ if elements of ranks i and j $(i < j)$ are ever compared.
Key insight: ranks i and j are compared iff **the first pivot chosen from within** $\{i, i+1, \ldots, j\}$ **is either** i **or** j. That happens with probability $\dfrac{2}{j-i+1}$.

$$\mathbb{E}[X] = \sum_{1 \le i < j \le n} \frac{2}{j-i+1} = \sum_{k=2}^{n} \frac{2}{k} \cdot (n-k+1) = 2(n+1)H_n - 4n.$$

So $\mathbb{E}[X] = \Theta(n \log n)$, and therefore expected runtime is $\Theta(n \log n)$.

Alternative View: Pivots Form a Random BST

Imagine inserting elements into a BST in the order in which they become pivots (the first pivot is the root, left-subarray pivots are its left subtree, ...). Two elements are compared exactly when one is an ancestor of the other — the same indicator argument, reinterpreted. The expected depth of the node ranked i is

$$H_i + H_{n-i+1} - 2 \le 2H_n = O(\log n),$$

so the expected total comparisons are $\sum_i O(\log n) = O(n \log n)$. Equivalently, arrange the pair-comparison probabilities in an upper-triangular matrix: every row is a truncated harmonic series of $O(\log n)$ "large" terms $(2/k)$, and all n rows together sum to $O(n \log n)$.

Quicksort Properties

- **Stable:** no (Lomuto and Hoare partitions swap non-adjacent elements).
- **In-place:** yes (only $O(\log n)$ recursion stack; no auxiliary array).
- **Practical speed:** usually faster than Mergesort or Heapsort despite the same

$\Theta(n \log n)$ bound, thanks to better cache locality and low constants.

Engineering Quicksort (What Real Libraries Do)

- **Insertion-sort cutoff:** for subarrays of size $\lesssim 16$, switch to insertion sort, which beats recursive quicksort on tiny nearly-sorted ranges.
- **Median-of-three pivot:** choose the median of first/middle/last elements — cheaply defeats already-sorted and reverse inputs.
- **Introsort:** switch to Heapsort after recursion depth exceeds $2\log_2 n$, retaining Quicksort's average speed with a guaranteed $O(n \log n)$ worst case (used by C++ `std::sort`).
- **3-way partition (Dutch national flag):** group equal keys together; without it, arrays with many duplicate keys degenerate toward $\Theta(n^2)$.

4.3 Randomized Selection (Quickselect)

Find the k-th smallest element of an unsorted array (median is the $k = \lceil n/2 \rceil$ case). Reuse Partition: after partitioning on a random pivot, the pivot is at its final rank $k = m - l + 1$ inside the current segment; only **one** recursive call is needed.

Algorithm 18 RandomizedSelect(A, l, r, k) — return the k-th smallest of $A[l..r]$

```
if l = r then
    return A[l]
end if
pick pivot uniformly at random, swap it to r
m ← PARTITION(A, l, r)
pos ← m − l + 1                                  ▷ rank of pivot within A[l..r]
if k = pos then
    return A[m]
else if k < pos then
    return RANDOMIZEDSELECT(A, l, m − 1, k)
else
    return RANDOMIZEDSELECT(A, m + 1, r, k − pos)                ▷ relative rank
end if
```

Worked example: find the 4-th smallest of $[7, 2, 9, 1, 4, 8, 3]$ (sorted: $1, 2, 3, \mathbf{4}, 7, 8, 9$).

Segment	k	pivot	result
$[7, 2, 9, 1, 4, 8, 3]$	4	3	partition $\to [2, 1, 3 \mid 7, 4, 8, 9]$, pivot rank $3 < 4$, go right
$[7, 4, 8, 9]$	1	9	rank $4 > 1$, go left
$[7, 4, 8]$	1	8	rank $3 > 1$, go left
$[7, 4]$	1	4	partition $\to [4 \mid 7]$, rank $1 = k$, **return** 4

Expected Time $\Theta(n)$

Call a pivot **good** when its rank lies in the middle half $[n/4, 3n/4]$; this happens with probability $1/2$. After a good pivot the surviving subproblem has size $\leq 3n/4$; bad pivots only add linear partitioning cost. Stages between good pivots are geometric with expected count 2, so the expected total work is

$$2cn + 2c\tfrac{3}{4}n + 2c\left(\tfrac{3}{4}\right)^2 n + \cdots = 2cn \sum_{i \geq 0} \left(\tfrac{3}{4}\right)^i = 8cn = \Theta(n).$$

Worst case remains $\Theta(n^2)$ (probability tiny); the deterministic **median-of-medians**

algorithm achieves worst-case $O(n)$ by splitting into groups of 5, whose pivot is guaranteed to have at least $3n/10$ elements on each side: $T(n) = T(n/5)+T(7n/10)+\Theta(n) = O(n)$.

Bonus: k-th of Two Sorted Arrays in $\Theta(\log k)$

Given sorted arrays A, B and rank k, compare $A[\lfloor k/2 \rfloor]$ and $B[\lfloor k/2 \rfloor]$ and discard the smaller half (none of its elements can be the k-th), then recurse with rank $k - \lfloor k/2 \rfloor$. Every step halves k; binary-search variants achieve $O(\log \min(m, n))$.

4.4 Hoare's Partition (Alternative)

Two pointers moving toward each other from opposite ends; swap inversions. Does slightly fewer swaps than Lomuto. Pivot is not guaranteed to end at index q; recur on $A[p..q]$ and $A[q+1..r]$, **skipping** q.

5 Heapsort

5.1 Priority Queue Motivation

A **priority queue** stores n elements, each with a priority (key), and supports `Insert` and `Extract-Max` (min version: negate priorities). Concrete uses: OS process scheduling (the `top` table — new/preempted processes are inserted; when the CPU is free, the lowest-PR process is extracted), event-driven simulation, Dijkstra's shortest paths, and streaming top-k.

Implementation	Insert	Extract-Max	Notes
Unsorted array/list	$O(1)$	$O(n)$	extract must scan everything
Sorted array/list	$O(n)$	$O(1)$	insert requires a shift
Balanced BST	$O(\log n)$	$O(\log n)$	correct bound but large pointer/balance overhead
Binary heap	$O(\log n)$	$O(\log n)$	compact array, no pointers, small constants

5.2 Heap Data Structure

A (binary) **max-heap** is a complete binary tree satisfying:

$$A[\text{Parent}(i)] \geq A[i] \quad \forall i.$$

(Min-heap: $\leq$ instead.)

Array Representation of a Heap

Given 1-indexed array $A[1..n]$:

$$\begin{aligned} \text{Parent}(i) &= \lfloor i/2 \rfloor, \\ \text{Left}(i) &= 2i, \\ \text{Right}(i) &= 2i+1. \end{aligned}$$

A complete binary tree on n nodes has height $\lfloor \lg n \rfloor$.

Height derivation. A complete tree of height h has all levels $0, \ldots, h-1$ full ($2^h - 1$ nodes) and at least one node on level h; hence

$$2^0 + 2^1 + \cdots + 2^{h-1} = 2^h - 1 < n \quad \Longrightarrow \quad n \leq 2^0 + 2^1 + \cdots + 2^h = 2^{h+1} - 1,$$

so $2^h \leq n < 2^{h+1}$ and $h = \lfloor \lg n \rfloor$. This $\Theta(\log n)$ height is why every heap operation is fast.

5.3 Max-Heapify (Bubble-Down)

Assume children of node i are heaps, but $A[i]$ may be too small. Fix by swapping $A[i]$ with the **larger** child, recursing down.

Algorithm 19 Max-Heapify

```
function MaxHeapify(A, i, heap_size)
    L ← 2i;    R ← 2i + 1
    largest ← i
    if L ≤ heap_size and A[L] > A[largest] then
        largest ← L
    end if
    if R ≤ heap_size and A[R] > A[largest] then
        largest ← R
    end if
    if largest ≠ i then
        swap A[i] ↔ A[largest]
        MaxHeapify(A, largest, heap_size)
    end if
end function
```

Time: $O(\log n)$ per call (proportional to tree height).

5.4 Build Max-Heap

Naive approach: run Max-Heapify top-down from the **last non-leaf** ($\lfloor n/2 \rfloor$) *down to* root 1. Leaves are trivially heaps of size 1.

Algorithm 20 Build Max-Heap

```
function BuildMaxHeap(A, n)
    for i = ⌊n/2⌋ downto 1 do
        MaxHeapify(A, i, n)
    end for
end function
```

Build-Max-Heap Cost is $O(n)$, Not $O(n \log n)$

Height-h nodes are $\leq n/2^{h+1}$; each MaxHeapify on a height-h node costs $O(h)$:

$$\sum_{h=0}^{\lfloor \lg n \rfloor} \frac{n}{2^{h+1}} \cdot O(h) = O\left(n \sum_{h=0}^{\infty} \frac{h}{2^h} \right) = O(2n) = O(n).$$

(Building from the leaves up is much cheaper than building from the root down.)

Per-level bookkeeping (take $n = 2^{h+1} - 1$, a perfectly full tree):

Level (root = 0)	# nodes	heapify cost per node	total for level
0	1	h	h
1	2	$h-1$	$2(h-1)$
2	4	$h-2$	$4(h-2)$
$\vdots$	$\vdots$	$\vdots$	$\vdots$
i	2^i	$h-i$	$2^i(h-i)$
h	2^h	0	0

Summing with $j = h - i$:

$$\sum_{i=0}^{h-1} 2^i (h-i) = 2^h \sum_{j=1}^{h} \frac{j}{2^j} = 2^h \left(2 - \frac{h+2}{2^h} \right) = 2^{h+1} - h - 2 < n,$$

using $\sum_{j=1}^{h} j/2^j = 2 - (h+2)/2^h$. Nearly all the work is at the bottom levels; the root level costs only h.

5.5 Heapsort Algorithm

(1) Build max-heap on $A[1..n]$. (2) Repeatedly: swap max (root $A[1]$) with last element, shrink heap, MaxHeapify root.

Algorithm 21 Heapsort

```
function HEAPSORT(A, n)
    BUILDMAXHEAP(A, n)
    for i = n downto 2 do
        swap A[1] ↔ A[i]
        heap_size ← i − 1
        MAXHEAPIFY(A, 1, heap_size)
    end for
end function
```

Heapsort Analysis

- Build: $O(n)$.
- $n-1$ swaps each followed by MaxHeapify: $(n-1)\,O(\log n) = O(n \log n)$.
- **Total:** $O(n \log n)$ (best = worst = average, unlike Quicksort).
- **Stable?** No.
- **In-place?** Yes (only $O(1)$ extra storage, no recursion stack if MaxHeapify is written iteratively).

5.6 Priority Queue Operations (Applications of Heaps)

Max-priority queue supports: `Maximum`, `ExtractMax`, `IncreaseKey`, `Insert`.

Algorithm 22 Priority Queue Operations

```
function EXTRACTMAX(A, heap_size)
    max ← A[1]
    A[1] ← A[heap_size];   heap_size ← heap_size − 1
    MAXHEAPIFY(A, 1, heap_size)
    return max
end function

function INCREASEKEY(A, i, key)
    if key < A[i] then error
    end if
    A[i] ← key
    while i > 1 and A[Parent(i)] < A[i] do                ▷ bubble up
        swap A[i] ↔ A[Parent(i)]
        i ← Parent(i)
    end while
end function

function INSERT(A, key, heap_size)
    heap_size ← heap_size + 1
    A[heap_size] ← −∞
    INCREASEKEY(A, heap_size, key)
end function
```

Each operation runs in $O(\log n)$ time.

Extract-Max trace on $[42, 29, 18, 14, 7, 6]$:

initial	save $m = 42$, move 6 to root	$[\mathbf{6}, 29, 18, 14, 7]$
bubble down	$6 < 29$, swap with larger child	$[\mathbf{29}, 6, 18, 14, 7]$
bubble down	$6 < \max(14, 7) = 14$, swap with 14	$[29, \mathbf{14}, 18, 6, 7]$

Returned 42; heap is now $[29, 14, 18, 6, 7]$.

Insert trace (bubble up): insert 20 into $[29, 14, 18, 6, 7]$:

append	$[29, 14, 18, 6, 7, \mathbf{20}]$	parent of 20 is 18
swap up	$20 > 18$	$[29, 14, \mathbf{20}, 6, 7, 18]$
stop	parent is $29 > 20$	done

5.7 Applications

5.7.1 Merging k Sorted Arrays

Combine k sorted arrays of length n into one sorted sequence of length kn.

Solution	Time	Idea
Sort everything	$O(kn \log(kn))$	ignore existing order
Divide & conquer	$O(kn \log k)$	pairwise 2-way merges, $\log_2 k$ rounds of kn work
Min-heap	$O(kn \log k)$	keep one current head per array in a heap of size k

Heap method: insert the first element of each array (tagged with its source array) into a min-heap; repeatedly extract-min and insert the next element from the same source array. Each of the kn elements costs one insert + one extract, each $O(\log k)$. Uses only $O(k)$ extra space, and it works for online/streaming inputs and huge files (external merge sort).

5.7.2 Delete and Update-Priority

Store `(key, priority, payload)` tuples, where `key` uniquely identifies the record.

- `Delete(key)`: locate its heap index k, swap $A[k]$ with $A[n]$, shrink, then **bubble up or bubble down** from k (either direction may be needed) — $O(\log n)$ once located.
- `UpdatePriority(key, p)`: set the new priority and bubble in the appropriate direction.

Locating an element by key requires a separate **position map** (hash table: key $\to$ heap index), updated on every swap. It preserves the $O(\log n)$ bounds but adds real overhead, which is why C++/Python standard heaps do not support it directly (`std::set` or lazy deletion are common workarounds).

5.7.3 The Skyline Problem

Given buildings (x_L, x_R, y), output the skyline key points (positions where the maximum height changes).

- **D&C solution:** recursively get the skylines of two halves; merge them like merge-sort while tracking current heights h_1, h_2 of each side. At each next boundary $x = \min(x_1, x_2)$, advance that side and append $(x, \max(h_1, h_2))$ only when it differs from the last emitted height. An $O(n)$ merge gives $T(n) = 2T(n/2) + \Theta(n) = O(n \log n)$.
- **Heap (sweep-line) solution:** sort all vertical boundaries $(x, y, \text{id}, L/R)$; sweep left to right, inserting a building at its left edge and deleting it at its right edge from a max-heap of live heights. The heap top is the current skyline height; emit a key point whenever it changes. $O(n \log n)$.

6 Linear-Time Sorting

All comparison-based sorts have *worst-case* $\Omega(n \log n)$ lower bound (decision-tree height). The following beat the bound by **not comparing keys directly**.

When Can We Beat $\Omega(n \log n)$?

We need extra information about the keys: their values lie in a **small integer range**, or they have structure that allows *digit-wise* processing.

6.1 The Decision-Tree Model and the Sorting Lower Bound

A **comparison-based** algorithm makes decisions only from pairwise comparisons ("$a_i \le a_j$?"); its behavior depends solely on the relative order, never on actual values. Merge, heap and quick sort are comparison-based; counting/radix/bucket sort are not.

Represent any execution as a **decision tree**: each internal node is one comparison with two outcomes (left/right branch); each leaf is one possible output. The **worst-case number of comparisons is the height** of the tree, and a binary tree with N leaves has height at least $\lceil \log_2 N \rceil$.

Warm-up: Lower Bound for Binary Search

Searching a sorted array of size n has $n+1$ distinct outputs — one per element, plus "not found." Hence the decision tree needs $\ge n+1$ leaves and height $\ge \log_2(n+1) = \Omega(\log n)$. Binary search matches this bound, so it is asymptotically optimal.

Exercise: The Defective Coin Balance

Among n coins, at most one is too heavy or too light (unknown which); a three-way balance reports left-heavy / balanced / right-heavy.

- **Lower bound:** there are $2n+1$ outcomes (n heavy, n light, all genuine), so the ternary decision tree has $\ge 2n+1$ leaves: height $\ge \log_3(2n+1)$ weighings.
- **Algorithm:** split into three equal groups C_1, C_2, C_3. Two weighings (C_1 vs C_2, then C_1 vs C_3) distinguish the 7 outcomes $C_1H, C_1L, C_2H, C_2L, C_3H, C_3L$, "none"; recurse into the defective group:

$$T(n) = T(n/3) + 2 = 2\log_3 n = O(\log n),$$

matching the lower bound up to constants.

Theorem: Sorting Requires $\Omega(n \log n)$ Comparisons

A sorting algorithm must be able to output any of the $n!$ input permutations, so its decision tree has $\ge n!$ leaves and height

$$h \ge \log_2(n!) = \sum_{i=1}^{n} \log_2 i \ge \sum_{i=\lceil n/2 \rceil}^{n} \log_2 i \ge \frac{n}{2} \log_2 \frac{n}{2} = \Omega(n \log n)$$

(equivalently use Stirling: $n! \ge (n/e)^n$, so $\log_2 n! = \Theta(n \log n)$). Therefore **every** comparison-based sort needs $\Omega(n \log n)$ comparisons in the worst case — Merge Sort and Heapsort are thus asymptotically optimal among comparison sorts. To do better we must leave the comparison model.

Algorithm vs problem lower bounds. "Insertion sort is $\Omega(n^2)$" is an algorithm bound (a particular algorithm's cost); "sorting is $\Omega(n \log n)$" is a problem bound (no comparison-based algorithm can beat it, proved independently of how clever the algorithm is). Problem lower bounds are the rarer and more valuable kind — they tell us when to

stop improving and point toward new models (non-comparison methods, randomization, average-case assumptions).

6.2 Counting Sort

Assume each of n input keys is an integer in $\{0, 1, \ldots, k\}$.

Algorithm 23 Counting Sort

```
function CountingSort(A, B, n, k)
    for i = 0 to k do
        C[i] ← 0
    end for
    for j = 1 to n do
        C[A[j]] ← C[A[j]] + 1                          ▷ count frequencies
    end for
    for i = 1 to k do
        C[i] ← C[i] + C[i − 1]                         ▷ prefix sums = "how many ≤ i"
    end for
    for j = n downto 1 do
        B[C[A[j]]] ← A[j]
        C[A[j]] ← C[A[j]] − 1
    end for
end function
```

Counting Sort Properties

- Runtime: $\Theta(n + k)$. If $k = O(n)$ then $\Theta(n)$ overall.
- **Stable:** yes (fill B *right-to-left* so equal keys preserve input order). This is critical for being used as a subroutine in Radix Sort.
- **In-place:** no (uses arrays C of size $k + 1$ and output B of size n).

Worked example: $A = [2, 5, 3, 0, 2, 3, 0, 3]$, $k = 5$.

i (key value)	0	1	2	3	4	5
frequencies C	2	0	2	3	0	1
prefix sums C	2	2	4	7	7	8

Placing right-to-left (equal keys keep their order): $B = [0, 0, 2, 2, 3, 3, 3, 5]$.

6.3 Radix Sort

Sort d-digit keys by applying a **stable** sort digit-by-digit, least significant digit first (LSD). Each digit sort is Counting Sort on that digit.

Algorithm 24 Radix Sort

```
function RadixSort(A, d, n)
    for i = 1 to d do
        Use a stable sort to sort A on digit i (the i-th least significant)
    end for
end function
```

Radix Sort Correctness & Complexity

- Why LSD works: after pass i, numbers are sorted correctly w.r.t. their low i digits;

next pass (higher digit) is stable so ties keep the lower-order ordering.
- Each digit in $\{0, \ldots, k-1\}$ sorted by Counting Sort in $\Theta(n+k)$.
- Total: $\Theta(d(n+k))$. If d is constant and $k = O(n)$, this is $\Theta(n)$.
- **Stable:** yes (because inner sort is stable).

Radix Sort Gotchas

- Must process digits **from LSD to MSD** (not the other way around).
- Inner sort **must be stable**; otherwise equal-digit elements lose their previously-computed order.

Correctness by Induction on Digit Position

Claim: after pass i, keys are sorted by their low-order i digits. Base $i = 0$ vacuous. Assume sorted by low $i-1$ digits; after stably sorting on digit i:
- two keys differing in digit i are correctly ordered by that more significant digit;
- two keys equal in digit i keep their pre-pass order by **stability**, which is correct ordering by the low $i-1$ digits (induction hypothesis).

Both cases give correctness by low i digits; after d passes the keys are fully sorted.

Worked example (3 LSD passes):

input	$[329, 457, 657, 839, 436, 720, 355]$
pass 1 (units)	$[720, 355, 436, 457, 657, 329, 839]$
pass 2 (tens)	$[720, 329, 436, 839, 355, 457, 657]$
pass 3 (hundreds)	$[329, 355, 436, 457, 657, 720, 839]$

Bit-Parameter Analysis (w-bit Integers)

Each "digit" spans $\log_2 k$ bits, so the number of passes is $d = \lceil w / \log k \rceil$ and total time is

$$\Theta\Big(\lceil \tfrac{w}{\log k} \rceil (n+k)\Big).$$

The $(n+k)$ factor is minimized at $k = n$ (equal split between pass count and bucket overhead), giving $\Theta(\lceil w/\log n \rceil n)$. In the standard word-RAM regime $w = O(\log n)$ (pointers/indices fit in a machine word; all machines from 8-bit to 64-bit satisfy this for realistic RAM sizes), radix sort runs in $\Theta(n)$. It also extends to fixed-length strings and to floating-point keys (their IEEE-754 bit order is monotone after a sign flip).

6.4 Bucket Sort

Assume input is drawn **uniformly at random** from $[0, 1)$. Distribute into n equal-sized buckets, sort each bucket (e.g. insertion sort), concatenate.

Algorithm 25 Bucket Sort

```
function BucketSort(A, n)
    Create B[0..n − 1] (each an empty list)
    for i = 1 to n do
        insert A[i] into B[⌊n · A[i]⌋]
    end for
    for i = 0 to n − 1 do
        sort list B[i] (e.g. insertion sort)
    end for
    Concatenate B[0], B[1], ..., B[n − 1] in order
end function
```

Bucket Sort Expected Time

With uniform input, each bucket has $O(1)$ elements in expectation; sorting each by insertion sort is $O(1)$ per bucket.

$$\mathbb{E}[T(n)] = \Theta(n) + \sum_{i=0}^{n-1} O(1) = \Theta(n).$$

Worst case (all elements fall in one bucket): $\Theta(n^2)$ if inner sort is insertion sort.

Why the Expected Cost Is Linear: $\mathbb{E}[X_0^2] = O(1)$

Let X_i be the number of elements in bucket i and $Y_j = \mathbf{1}\{A[j] \text{ lands in bucket } 0\}$, so $X_0 = \sum_{j=1}^{n} Y_j$. Insertion sort on bucket i costs $\Theta(X_i^2)$; by identical distributions,

$$\mathbb{E}\left[\sum_{i=0}^{n-1} X_i^2\right] = n\,\mathbb{E}[X_0^2].$$

Expand the square and use independence of distinct Y_j ($\Pr[Y_j = 1] = 1/n$):

$$\mathbb{E}[X_0^2] = \sum_j \mathbb{E}[Y_j^2] + 2\sum_{j<\ell} \mathbb{E}[Y_j Y_\ell] = n \cdot \frac{1}{n} + 2\binom{n}{2}\frac{1}{n^2} = 1 + 1 - \frac{1}{n} = 2 - \frac{1}{n} = O(1).$$

So bucket sorting costs $O(n)$ expected time total (plus the $O(n)$ distribution/concatenation passes). If inputs are not in $[0, 1)$, normalize first; with non-uniform distributions a custom bucket partition is needed.

Sorting Algorithm Cheat Sheet

Algorithm	Worst	Average	Space	Comparison?	Stable	Cache	Parallel
Insertion	$\Theta(n^2)$	$\Theta(n^2)$	$\Theta(1)$	yes	yes	good	no
Mergesort	$\Theta(n \log n)$	$\Theta(n \log n)$	$\Theta(n)$	yes	yes	good	good
Quicksort	$\Theta(n^2)$	$\Theta(n \log n)$	$\Theta(\log n)$	yes	no	good	good
Heapsort	$\Theta(n \log n)$	$\Theta(n \log n)$	$\Theta(1)$	yes	no	bad	no
Counting	$\Theta(n + k)$	$\Theta(n + k)$	$\Theta(n + k)$	no	yes	bad	no
Radix	$\Theta(\lceil w/\log n \rceil n)$	same	$\Theta(n)$	no	yes	bad	good
Bucket	$\Theta(n^2)$	$\Theta(n)$ expected	$\Theta(n)$	no	yes	bad	good

7 Greedy Algorithms

A **greedy algorithm** always makes the choice that looks best at the moment and adds it to the current partial solution. It builds a solution to an optimization problem through a sequence of choices that are:

- **Feasible:** satisfying the given constraints.
- **Locally optimal:** make the choice that looks best right now, without considering the future.
- **Irrevocable:** once a choice is made it is never undone (no backtracking).

When Does Greedy Work?

- A greedy algorithm is correct when the problem has the **greedy-choice property** (a locally optimal choice can be part of some global optimum) together with **optimal substructure** (an optimal solution contains optimal solutions to subproblems).
- When they work, greedy algorithms are usually the simplest and most efficient algorithms available. A sorting step often dominates the cost, giving $\Theta(n \log n)$.
- When they fail, they may still give fast **approximations**.

Greedy Correctness Must Be Proven!

"Picking the locally best option" is not automatically optimal. The correctness (optimality) of a greedy algorithm is usually non-obvious. Two standard proof techniques:

1. **Exchange argument:** transform any optimal solution into the greedy solution one choice at a time, without worsening it (Interval Scheduling, Fractional Knapsack, Huffman).
2. **Greedy stays ahead / match a lower bound:** show the greedy solution reaches a universal lower bound (Interval Partitioning vs. depth).

7.1 Warm-up: Change-Making

Input: amount n and coin denominations $d_1 > d_2 > \cdots > d_m$. Output: the minimum number of coins summing to n.

Greedy rule: give as many coins as possible of the largest denomination, then the second largest, and so on.

Example (Hong Kong coins): $n = 19$, denominations $10, 5, 2, 1$.

$$19 = 10 + 5 + 2 + 2 \quad \Longrightarrow \quad 4 \text{ coins, and this is optimal.}$$

Greedy Change-Making Is Not Always Optimal

British coinage before 1971 used denominations $1, 3, 6, 12, 24, 30, 60$. For $n = 48$:

- Greedy: $48 = 30 + 12 + 6$ (3 coins).
- Optimal: $48 = 24 + 24$ (2 coins).

The greedy choice (take 30) "steals" capacity that two better coins could have used. For arbitrary denominations, change-making requires **dynamic programming**.

7.2 Warm-up: Max Sum of Non-Adjacent Elements

Given n positive numbers $A[1], \ldots, A[n]$, pick a subset of **pairwise non-adjacent** elements with maximum total sum. Example: $A = (1, 8, 6, 3, 6) \Rightarrow$ pick $A[2], A[5]$ with sum $8 + 6 = 14$.

Greedy attempt: repeatedly take the largest remaining element, then delete it and its two neighbors.

Algorithm 26 Greedy Non-Adjacent Selection (Not Optimal!)

$S \leftarrow \emptyset$
while some elements remain in A **do**
  pick index i with the largest remaining $A[i]$
  $S \leftarrow S \cup \{A[i]\}$
  mark $A[i-1], A[i], A[i+1]$ as deleted
end while
return S

Counterexample: $(1, 4, 6, 4)$

- Greedy takes the maximum 6, which forces deleting both neighbors 4 and 4, then takes 1: sum $= 6 + 1 = 7$.
- Optimal: take the two 4's (positions 2 and 4): sum $= 8 > 7$.

The optimal solution is given by DP: $\mathrm{OPT}(i) = \max\big(\mathrm{OPT}(i-1),\ \mathrm{OPT}(i-2) + A[i]\big)$.

7.3 Interval Scheduling

Job j starts at s_j and finishes at f_j. Two jobs are **compatible** if their intervals do not overlap. Goal: select a **maximum-size** subset of mutually compatible jobs.

7.3.1 Three Tempting Greedy Rules (All Fail)

Counterexamples

1. **Earliest start time.** One long job $e = [0, 10)$ starting first can block four short jobs $[1, 2), [2, 3), [3, 4), [4, 5)$: greedy chooses $\{e\}$ (size 1) instead of $\{a, b, c, d\}$ (size 4).
2. **Shortest interval.** Let $a = [0, 4)$, $b = [4, 8)$, and a very short $c = [3.5, 4.5)$ overlapping both. Shortest-first chooses $\{c\}$ (size 1) instead of $\{a, b\}$ (size 2).
3. **Fewest conflicts.** Count each job's overlapping jobs c_j and schedule in ascending c_j. A job conflicting with only a few others can still sit in the middle of a long compatible chain and, once chosen, force several of them out — the slide example forces $\{a, f, d\}$ (size 3) instead of $\{a, b, c, d\}$ (size 4). The number of conflicts is a global statistic and does not measure local opportunity cost.

7.3.2 The Correct Rule: Earliest Finish Time

Greedy algorithm: consider jobs in increasing order of finish time; take a job whenever it is compatible with the jobs already chosen. Intuition: finishing early **leaves the maximum remaining time** for all later jobs.

Algorithm 27 Interval Scheduling (Earliest Finish Time)

```
Sort jobs so that f_1 ≤ f_2 ≤ ··· ≤ f_n
A ← ∅;   last ← 0                          ▷ finish time of the last selected job
for j = 1 to n do
    if s_j ≥ last then                     ▷ compatible: ends are inclusive, so ≥ is fine
        A ← A ∪ {j};   last ← f_j
    end if
end for
return A
```

Running time: $\Theta(n \log n)$, dominated by sorting; the scan itself is $\Theta(n)$.

Worked trace. Jobs sorted by finish time:

$$j_1 = [0, 2),\ j_2 = [1, 4),\ j_3 = [3, 5),\ j_4 = [4, 6),\ j_5 = [5, 7).$$

- Select j_1 (earliest finish), set $last = 2$.
- Skip j_2: $s_2 = 1 < 2$, overlaps j_1.
- Select j_3: $s_3 = 3 \geq 2$, set $last = 5$.
- Skip j_4: $s_4 = 4 < 5$.
- Select j_5: $s_5 = 5 \geq 5$.

Greedy returns $\{j_1, j_3, j_5\}$, size 3 — the maximum possible.

7.3.3 Correctness Proof (Exchange Argument)

Theorem: Earliest-Finish-Time Greedy Is Optimal

Let $G = (i_1, i_2, \ldots, i_k)$ be the jobs selected by greedy and $O = (j_1, j_2, \ldots, j_m)$ be any optimal schedule, both in selection order.

1. Suppose $G \neq O$. Find the largest r such that $i_1 = j_1, \ldots, i_r = j_r$ (the common prefix; $r = 0$ if they differ immediately).
2. By the greedy rule, i_{r+1} is the compatible job with the **earliest finish**, so $f_{i_{r+1}} \leq f_{j_{r+1}}$.
3. Build O^* from O by replacing j_{r+1} with i_{r+1}. Because i_{r+1} finishes no later than j_{r+1}, it is still compatible with $j_{r+2}, \ldots, j_m$; hence O^* is a feasible schedule with the same size as O, so O^* is also optimal.
4. O^* now shares the first $r + 1$ jobs with G. Repeat the replacement; the common prefix grows each time, so eventually O is transformed into G with its size (value) unchanged.

Therefore $|G| = |O|$, i.e. greedy is optimal. Equivalently, the greedy schedule "stays ahead": the r-th greedy job finishes no later than the r-th job of any schedule, so greedy can always fit at least as many jobs.

7.4 Exercise: Interval Covering

Given n intervals that together cover $[0, 1]$; interval x has start $x.s$ and finish $x.f$. Find the **minimum** number of intervals whose union covers $[0, 1]$. (Assume all endpoints are distinct.)

Greedy idea: maintain the rightmost covered point x. Among all intervals covering x ($I.s \leq x < I.f$), pick the one extending **furthest to the right**; advance x to its finish and repeat until $x \geq 1$.

Algorithm 28 Greedy Interval Covering

$x \leftarrow 0$; $G \leftarrow \emptyset$
while $x < 1$ **do**
  pick interval I with $I.s \le x$ that maximizes $I.f$
  $G \leftarrow G \cup \{I\}$; $x \leftarrow I.f$
end while
return G

Worked trace. Intervals $A = [0, 0.3]$, $B = [0.2, 0.6]$, $C = [0.5, 0.8]$, $D = [0.75, 1]$:

- $x = 0$: only A covers 0 $\Rightarrow$ take A, $x = 0.3$.
- $x = 0.3$: B covers 0.3 and extends furthest $\Rightarrow$ take B, $x = 0.6$.
- $x = 0.6$: C covers 0.6 (D starts at 0.75) $\Rightarrow$ take C, $x = 0.8$.
- $x = 0.8$: take D, $x = 1$. Done with 4 intervals.

Optimality Sketch & Implementation

Let $O = (o_1, \ldots, o_m)$ be optimal and $G = (g_1, \ldots, g_k)$ greedy, ordered by start time. Let i be the first index with $o_i \neq g_i$. The greedy choice extends at least as far: $g_i.f \ge o_i.f$. Replacing o_i by g_i keeps the cover valid; iterating converts O into G, so $k \le m$ and greedy is optimal. Sorting plus a single sweep (among intervals already started, track the maximum finish) gives $O(n \log n)$.

7.5 Fractional Knapsack

Input: n items, item i has weight w_i and value v_i; knapsack capacity W. Choose fractions $0 \le x_i \le 1$ maximizing $\sum_i x_i v_i$ subject to $\sum_i x_i w_i \le W$.

- **0/1 knapsack:** $x_i \in \{0, 1\}$ (take the whole item or nothing) — greedy fails; needs DP.
- **Fractional knapsack:** $x_i \in [0, 1]$ (items are divisible, e.g. liquids, grain) — greedy is optimal.

Greedy rule: sort items by **value per unit weight** v_i/w_i decreasing; take as much as possible of each item in that order. Running time $\Theta(n \log n)$.

Algorithm 29 Fractional Knapsack (Greedy)

Sort items so that $\frac{v_1}{w_1} \ge \frac{v_2}{w_2} \ge \cdots \ge \frac{v_n}{w_n}$
$w \leftarrow W$ ▷ remaining capacity
for $i = 1$ **to** n **do**
  if $w_i \le w$ **then**
    $x_i \leftarrow 1$; $w \leftarrow w - w_i$
  else
    $x_i \leftarrow w/w_i$; $w \leftarrow 0$; **return**
  end if
end for

Worked example. $W = 50$; items $(w_i, v_i) = (10, 60), (20, 100), (30, 120)$, with ratios $6, 5, 4$.

1. Take all of item 1: used weight 10, value 60.
2. Take all of item 2: used weight 30, value 160.
3. Fill the remaining 20 units with a fraction $20/30 = 2/3$ of item 3: added value $\frac{2}{3} \cdot 120 = 80$.

Optimal total: $60 + 100 + 80 = \mathbf{240}$, weight exactly 50.

Why the Same Greedy Fails for 0/1 Knapsack

On the example above, "take by ratio while the whole item fits" grabs items 1 and 2 (weight 30, value 160) and then item 3 does not fit — leaving 20 units of capacity wasted. The 0/1 optimum takes items 2 and 3: weight 50, value 220 $>$ 160. Being forced to leave capacity unused is exactly what divisibility (the fractional version) avoids.

Correctness Proof (Exchange Argument)

Assume $\sum_i w_i \geq W$ (otherwise take everything; trivially optimal), so both greedy $G = (x_1, \ldots, x_k, 0, \ldots, 0)$ (where $x_i = 1$ for $i < k$ and $0 \leq x_k \leq 1$) and any optimum $O = (y_1, \ldots, y_n)$ fill the knapsack exactly.

1. Let i be the first index where $x_i \neq y_i$. By the greedy ordering and the prefix agreement, greedy takes at least as much of the best remaining item: $x_i \geq y_i$.
2. Increase y_i to x_i and compensate by removing total weight $(x_i - y_i)w_i$ from items $i+1, \ldots, n$ (enough such weight exists because both solutions pack exactly W).
3. Value cannot decrease: every item $j > i$ has ratio $v_j/w_j \leq v_i/w_i$, so swapping weight onto item i replaces "cheaper" weight with "dearer" weight.
4. Since O was optimal, its value also cannot increase, so it stays exactly the same: the modified O is still optimal and now agrees with G up to index $i+1$.

Repeating converts O into G with unchanged value $\Rightarrow$ greedy is optimal.

7.6 Interval Partitioning

Lecture j occupies $[s_j, f_j)$. Goal: schedule all lectures using the **minimum number of classrooms** so that no two lectures overlap in the same room.

Greedy algorithm: sort lectures by increasing **start** time; assign each lecture to any classroom whose last lecture has finished; if none exists, open a new classroom.

Algorithm 30 Interval Partitioning (Greedy by Start Time)

```
Sort lectures so that s_1 ≤ s_2 ≤ ··· ≤ s_n
H ← empty min-heap of classroom finish times
for j = 1 to n do
    if H ≠ ∅ and min(H) ≤ s_j then
        ExtractMin(H)                          ▷ reuse that classroom
    end if
    Insert(H, f_j)                             ▷ new room if none was reused
end for
return |H|                                     ▷ number of classrooms used
```

Worked trace. Lectures (sorted by start): $A = [1, 4)$, $B = [2, 6)$, $F = [3, 3.5)$, $C = [4, 5)$, $D = [5, 7)$, $E = [6, 8)$.

Lecture	Decision	Heap of finish times	Overlapping rooms
$A[1, 4)$	open room 1	$\{4\}$	1
$B[2, 6)$	$4 > 2$, open room 2	$\{4, 6\}$	2
$F[3, 3.5)$	$4 > 3$, open room 3	$\{3.5, 4, 6\}$	3
$C[4, 5)$	$3.5 \leq 4$, reuse	$\{4, 5, 6\}$	3
$D[5, 7)$	$4 \leq 5$, reuse	$\{5, 6, 7\}$	3
$E[6, 8)$	$5 \leq 6$, reuse	$\{6, 7, 8\}$	3

Three rooms are used; at time 3 lectures A, B, F overlap simultaneously, so 3 is unavoidable.

Depth Lower Bound and Correctness

The **depth** of a set of intervals is the maximum number of intervals covering any single point in time.

- **Lower bound:** every schedule needs at least depth classrooms — at that instant all those lectures are running and each needs its own room.
- **Greedy matches it:** suppose greedy opens its d-th classroom for lecture j. At that moment all $d-1$ existing rooms have finish time $> s_j$; by the start-time ordering their lectures started at or before s_j, so together with j there are d lectures covering time s_j. Hence depth $\geq d$.

Greedy uses exactly (no more than) the depth $\Rightarrow$ optimal.

Running Time

- Naively scanning all rooms for compatibility: $O(n^2)$.
- Key observation: if j is incompatible with the room whose last lecture finishes **earliest**, it is incompatible with every room. So one min-heap keyed by room finish times suffices: each lecture does one `Min` check ($O(1)$) and at most one `ExtractMin`/`Insert` ($O(\log n)$ each).
- Total: $\Theta(n \log n)$ including the sort.

Sorting by Finish Time with Arbitrary Room Assignment Fails

Intervals $(1,3),(2,4),(4,5),(3,6)$, processed by increasing finish time:

1. $(1,3) \rightarrow$ room 1 (finishes 3).
2. $(2,4)$: no room free $\rightarrow$ room 2 (finishes 4).
3. $(4,5)$: both rooms are free; **arbitrarily** choose room 1 (finish 3) $\rightarrow$ room 1 now finishes 5.
4. $(3,6)$: neither room is free $(5, 4 > 3) \rightarrow$ room 3.

This uses 3 rooms although depth is 2 (optimal: room 1 gets $(1,3),(3,6)$; room 2 gets $(2,4),(4,5)$). **Fix (tie-breaking):** if several rooms are available, reuse the one with the **latest** finish time (best fit). Implement with a BST of rooms ordered by finish time: take the predecessor with finish $\leq s_j$ and update it to f_j; total $O(n \log n)$.

8 Huffman Coding

Huffman coding is another greedy algorithm: it builds an optimal **prefix-free binary code** for a file whose character frequencies are known.

8.1 Motivation: Variable-Length Encoding

There are over 10^6 characters across all languages: a fixed-length code would need ≥ 21 bits per character, while English needs only 7 (ASCII). Internet traffic is predominantly English, so **variable-length** codes — short codewords for frequent characters, long codewords for rare ones — save many bits in expectation.

8.2 Prefix-Free Codes

A code is a **prefix-free code** ("prefix code") if no codeword is a prefix of any other codeword.

Prefix-Free ⇒ Unique Decodability

Theorem. Every message encoded by a prefix-free code is uniquely decodable: reading left to right, the bits can match only one codeword (it cannot be a prefix of a longer one), so peel it off and continue.

Examples.

- $C_1 = \{a = 00, b = 01, c = 10, d = 11\}$ (fixed-length): $010011 = 01\,|\,00\,|\,11 \rightarrow$ "bad".
- $C_2 = \{a = 0, b = 110, c = 10, d = 111\}$ (prefix-free): $1100111 = 110\,|\,0\,|\,111 \rightarrow$ "bad".
- $C_3 = \{a = 1, b = 110, c = 10, d = 111\}$ (not prefix-free, since 1 is a prefix of 110 and 111): 1101111 can be parsed as $110\,|\,1\,|\,111$ = "bad" or $1\,|\,10\,|\,1\,|\,111$ = "acad" — ambiguous!

Why Prefix Codes in Practice?

UTF-8 marks continuation bytes with the prefix `10`. If bytes are corrupted or dropped, the decoder self-synchronizes: skip ahead to the next byte not beginning with `10` rather than losing the rest of the file.

8.3 The Encoding Problem

Given characters with known frequencies, find a prefix code minimizing the total encoded length (compression).

Example file: 100,000 characters over $\{a, b, c, d, e, f\}$ with frequencies in thousands:

	a	b	c	d	e	f	Total (bits)
Frequency ($\times 1000$)	45	13	12	16	9	5	100,000 chars
Fixed-length code	000	001	010	011	100	101	$3 \cdot 100{,}000 = 300{,}000$
Variable-length code	0	101	100	111	1101	1100	224,000

The variable code costs

$(45{\cdot}1{+}13{\cdot}3{+}12{\cdot}3{+}16{\cdot}3{+}9{\cdot}4{+}5{\cdot}4){\cdot}1000 = (45{+}39{+}36{+}48{+}36{+}20){\cdot}1000 = 224{,}000$ bits,

a 25% reduction.

8.4 Tree Correspondence and Weighted External Path Length

Prefix Codes ↔ Binary Trees

Represent a prefix code by a binary tree: label the left edge 0 and the right edge 1; the codeword of a character is the bit string on the root-to-leaf path.

- Every **leaf** stores one character; an internal node has two children (a "full" tree — no leaf codeword can be a prefix of another).
- Leaf depth $d(a_i)$ equals codeword length.
- Set leaf weight to the character frequency $f(a_i)$. The **weighted external path length**

$$B(T) = \sum_{i=1}^{n} f(a_i)\, d(a_i)$$

is exactly the total encoded length. Finding the minimum-cost prefix code = finding a minimum-weighted-external-path-length binary tree!

In the example, the fixed-length tree (depth 3 for all six leaves) costs $(45{+}13{+}12{+}16{+}9{+}5){\cdot}3 = 300$ (thousand bits); the variable tree costs $45{\cdot}1{+}(12{+}13{+}16){\cdot}3{+}(9{+}5){\cdot}4 = 224$.

8.5 Huffman's Greedy Algorithm

Greedy idea (a "merge"):

1. Pick the two characters x, y with the **smallest** weights.
2. Join them as sibling leaves under a new internal node z with weight $f(z) = f(x){+}f(y)$.
3. Remove x, y from the alphabet and insert z.
4. Repeat until one node (the root) remains.

Algorithm 31 Huffman

```
function HUFFMAN(A)
    build a min-priority queue Q on A, keyed by weight
    for i = 1 to n − 1 do
        allocate a new node z
        x ← EXTRACTMIN(Q)
        y ← EXTRACTMIN(Q)
        z.left ← x;    z.right ← y
        z.weight ← x.weight + y.weight
        INSERT(Q, z)
    end for
    return EXTRACTMIN(Q)                    ▷ root of the Huffman tree
end function
```

Running time: $O(n \log n)$ with a binary min-heap ($n-1$ merge rounds, each with two extract-mins and one insert, each $O(\log n)$).

8.6 Worked Example

Initial alphabet (weights): f:5, e:9, c:12, b:13, d:16, a:45.

Round	Merge	Priority queue afterwards
0	—	$\{5f,\ 9e,\ 12c,\ 13b,\ 16d,\ 45a\}$
1	$5f + 9e = 14$	$\{12c,\ 13b,\ 14,\ 16d,\ 45a\}$
2	$12c + 13b = 25$	$\{14,\ 16d,\ 25,\ 45a\}$
3	$14 + 16d = 30$	$\{25,\ 30,\ 45a\}$
4	$25 + 30 = 55$	$\{45a,\ 55\}$
5	$45a + 55 = 100$	$\{100\}$ (root)

Reading root-to-leaf paths (left = 0, right = 1) gives:

Character	a	c	b	f	e	d
Codeword	0	100	101	1100	1101	111
Length	1	3	3	4	4	3

Two Ways to Compute the Total Cost

- By leaf depths: $B(T) = \sum_i f_i d_i = 45{\cdot}1 + (12 + 13 + 16){\cdot}3 + (9 + 5){\cdot}4 = 224$.
- By merge weights: $B(T) = 14 + 25 + 30 + 55 + 100 = 224$.

The sums agree because a leaf of weight f contributes f to the merge weight of every internal node on its path to the root — exactly one contribution per level of depth. This identity is the key to both the proof and the optimal-merge application.

8.7 Correctness Proof

Lemma 1: An Optimal Tree Is Full

Every internal node in an optimal prefix-code tree has exactly two children. If an internal node had only one child, contract that node (replace it by its child): no leaf gets deeper, and some leaves get shallower, strictly decreasing $B(T)$ — contradiction.

Lemma 2: Swapping Leaves

Let T' be obtained from T by swapping two leaves x and b. If $f(x) \le f(b)$ and $d(x) \le d(b)$ (move the smaller-frequency leaf down and the larger one up), then

$$B(T') = B(T) + \big(f(x) - f(b)\big)\big(d(b) - d(x)\big) \le B(T),$$

since the first factor is ≤ 0 and the second is ≥ 0. So placing frequent characters nearer the root never increases the cost.

Lemma 3: The Two Smallest Weights Can Be Siblings at Maximum Depth

Let x, y be the two least frequent characters. Take any optimal tree T (full by Lemma 1); it contains two sibling leaves b, c at the deepest level. Assuming $f(x) \le f(b)$ and $f(y) \le f(c)$ (relabel if needed), swap $x \leftrightarrow b$ and $y \leftrightarrow c$. By Lemma 2 the cost does not increase; since T was optimal, the resulting tree is also optimal and now has x, y as sibling leaves at the deepest level.

Lemma 4: Contracting Two Sibling Leaves

If x, y are sibling leaves in T, contract them into one leaf z with $f(z) = f(x) + f(y)$ to obtain T', then

$$\begin{aligned} B(T) &= B(T') - f(z)d(z) + f(x)(d(z) + 1) + f(y)(d(z) + 1) \\ &= B(T') + f(x) + f(y). \end{aligned}$$

Theorem: Huffman Produces an Optimal Tree (Induction on n)

Base ($n = 2$): the only full tree has two leaves at depth 1; it is optimal.
Inductive step: assume Huffman is optimal for alphabets of size $n - 1$. For size n, let H be the Huffman tree and x, y its first merge (the two smallest weights, siblings in H). Contract them to z with $f(z) = f(x) + f(y)$: the contracted tree H' is exactly the tree Huffman builds on the $n - 1$-character alphabet $A' = A - \{x, y\} \cup \{z\}$, so by the induction hypothesis H' is optimal for A', and by Lemma 4

$$B(H) = B(H') + f(x) + f(y).$$

By Lemma 3 there is an optimal tree T for A with x, y as siblings; contract them to get T' for A'. Then

$$B(H) = B(H') + f(x) + f(y) \leq B(T') + f(x) + f(y) = B(T).$$

Since T is optimal, H is optimal too. □

Huffman Caveats

- Huffman is optimal **among prefix codes**, not among all possible codes (arithmetic coding can exploit fractional code lengths).
- Each merge must join the two current smallest weights — joining the smallest with an arbitrary node is not optimal.
- Ties in weight may be broken arbitrarily: different trees can result, all with the same minimum cost.
- Edge case $n = 1$: the single character gets the empty codeword (conventionally one bit may be used).

8.8 Application: Optimal Two-Way Merge

Given n sorted lists $L_1, \ldots, L_n$, merge them into one sorted list, merging only two lists at a time; merging lists of sizes p and q costs $p + q$ comparisons and produces a list of size $p + q$. Find the merge pattern minimizing total comparisons.

A merge pattern is a binary tree: leaves are the initial lists (weight = list size), each internal node is a merge, and the total cost equals the **weighted external path length** — so Huffman (always merge the two smallest lists) is optimal.

Example: three lists of sizes $30, 20, 10$:

Pattern	Computation	Total comparisons
L_1, L_2 first, then L_3	$30 + 20 = 50$, then $50 + 10 = 60$	$50 + 60 = 110$
L_1, L_3 first, then L_2	$30 + 10 = 40$, then $40 + 20 = 60$	$40 + 60 = 100$
L_2, L_3 **first, then** L_1 **(Huffman)**	$20 + 10 = 30$, then $30 + 30 = 60$	$\mathbf{30 + 60 = 90}$

Algorithm 32 Optimal Binary Merge Tree

```
build a min-heap from the n initial list sizes
while heap size ≥ 2 do
    a ← ExtractMin;   b ← ExtractMin
    create an internal node of size a + b with children a, b
    Insert(a + b)
end while
```

Time $O(n \log n)$.

Worked example: initial sizes $2, 3, 5, 7, 9$:

Round	Merge	Heap afterwards
1	$2+3=5$	$\{5,5,7,9\}$
2	$5+5=10$	$\{7,9,10\}$
3	$7+9=16$	$\{10,16\}$
4	$10+16=26$	$\{26\}$

Total cost (sum of merge sizes): $5+10+16+26 = \mathbf{57}$, which equals the weighted external path length of the merge tree:

$$2\cdot 3+3\cdot 3+5\cdot 2+7\cdot 2+9\cdot 2 = 6+9+10+14+18 = 57.$$

Greedy Algorithms at a Glance

Problem	Greedy choice	Complexity	Correctness idea
Interval scheduling	earliest finish time	$\Theta(n \log n)$	exchange argument
Interval covering	furthest reach covering frontier	$O(n \log n)$	exchange argument
Fractional knapsack	largest value/weight	$\Theta(n \log n)$	exchange argument
Interval partitioning	earliest start, reuse a room	$\Theta(n \log n)$	match depth lower bound
Huffman coding	merge two smallest weights	$O(n \log n)$	exchange + induction
Optimal merge	merge two smallest lists	$O(n \log n)$	same as Huffman